

Sparse Feature Analysis of Deep Layer Expansion: A Mechanistic Interpretation via SAE

Authors: Jin Yanyan, Zhao Lei

Discussion Paper v3 — Extends Zhao (2026) with Sparse Autoencoder analysis

Abstract

Zhao (2026) demonstrated that expert-level prompts induce “Deep Layer Expansion”—a 60-100% increase in Effective Intrinsic Dimension (EID) at deep layers. However, EID is a global metric that does not reveal **which semantic features** are activated. In this paper, we apply Sparse Autoencoder (SAE) analysis to decompose the activation differences between prompt styles. Using Goodfire’s Llama-3.3-70B SAE (Layer 50, 65,536 features), we find that: (1) “Explain to a novice” activates **17% more features** than “explain to an expert” (132.4 vs 113.1 on average); (2) **369 features are exclusively activated by novice prompts** vs 208 for expert prompts; (3) **10 features show perfect separation** between Novice vs Expert conditions; (4) Through AutolInterp analysis (6 conditions \times 50 topics = 300 samples), we discover these features exhibit **semantic subdivision**—encoding distinct dimensions such as “expert identity,” “serious attitude,” “depth requirement,” and “technical analysis”; (5) **UMAP visualization confirms that 6 prompt conditions form distinct clusters** in both raw activation space and SAE feature space, with SAE acting as a semantic denoiser that merges noise-only conditions (standard/padding/spaces) while preserving semantic distinctions (novice/expert/guru). These findings suggest prompt effects are **compositional**, with different elements triggering different feature subsets.

Keywords: Sparse Autoencoder, Interpretability, Prompt Engineering, Feature Activation, Llama

1. Introduction

1.1 Background: The EID Puzzle

Zhao (2026) established a striking empirical finding: expert-level prompts increase deep-layer EID by 60-100% compared to standard prompts. The **Manifold Teleportation** hypothesis explains this as expert signals navigating activation trajectories toward high-dimensional semantic regions.

However, a fundamental question remains unanswered:

What exactly changes inside the model when EID increases?

EID is computed from singular value entropy—a global summary statistic. It tells us the representation is “higher dimensional,” but not **which dimensions** are activated. Two representations with identical EID could involve completely different semantic features.

1.2 SAE: A Window into Sparse Features

Sparse Autoencoders (SAE) provide a tool to decompose dense activations into interpretable features. The core idea:

Hidden State [8192] → SAE Encoder → Sparse Features [65536]
(most entries = 0)

Each of the 65,536 features (ideally) corresponds to a distinct semantic concept. The sparsity constraint ensures that only a small subset (~100-200) are active for any given input.

Key insight: If expert prompts induce higher EID, SAE analysis can reveal whether this reflects:
- (A) More features being activated (activation count \uparrow) - (B) Different features being activated (feature set changes) - (C) Stronger activation of the same features (activation intensity \uparrow)

1.3 Our Contribution

We apply SAE analysis to the prompt comparison paradigm from Zhao (2026), with a twist: instead of “standard vs expert,” we compare “**novice vs expert**”—two prompts that should induce opposite effects on explanation complexity.

Hypothesis: If “expert” prompts induce dimensional expansion through specialized knowledge activation, then “novice” prompts should induce even greater expansion—because explaining to a beginner requires activating **more** semantic units (background knowledge, analogies, simplified models).

2. Methods

2.1 SAE Model

We use Goodfire’s publicly released SAE for Llama-3.3-70B-Instruct: - **Layer:** 50 (of 80) - **Input dimension:** 8,192 (Llama’s hidden size) - **Feature dimension:** 65,536 - **Architecture:** Linear encoder/decoder with ReLU activation

SAE encoding: $f = \text{ReLU}(x \cdot W_{enc}^T + b_{enc})$

2.2 Prompt Conditions

We test six prompt styles on 50 technical topics:

Condition	Template
standard padding	“Please explain {topic}.” “Please explain {topic}.” + filler text (length control)
spaces	“Please explain {topic}.” + whitespace (length control)
novice	“Please explain {topic} to a complete novice.”
expert	“Please explain {topic} to a domain expert.”
guru	“As {famous_name}, explain {topic}.”

2.3 Measurement Protocol

1. Process each prompt through Llama-3.3-70B-Instruct
2. Extract Layer 50 hidden state at the last token position
3. Apply SAE encoder to obtain 65,536-dimensional sparse representation
4. Compare activation patterns across conditions

2.4 Metrics

- **Activation count:** Number of features with value > 0
 - **Activation frequency:** % of samples where a feature is active
 - **Exclusive features:** Features active in only one condition
 - **Perfect separators:** Features with 100% activation in one condition, 0% in the other
-

3. Results

3.1 Activation Count: Novice > Expert

Condition	Avg. Active Features	Max Activation
novice	132.4	4.71
standard	112.4	4.11
padding	126.0	6.11
guru	115.0	4.39
expert	113.1	5.31
spaces	99.0	4.80

Key finding: Novice prompts activate **17% more features** than expert prompts (132.4 vs 113.1).

This confirms Hypothesis (A): EID differences reflect **more features being activated**, not just stronger activation of the same features.

3.2 Exclusive Features: Asymmetric Activation

Metric	Novice	Expert
Exclusive features	369	208
Ratio	1.77x	1.00x

369 features are activated only by novice prompts, compared to 208 for expert prompts—a 77% asymmetry.

This confirms Hypothesis (B): Different prompt styles activate **different feature sets**, not merely different intensities of the same features.

3.3 Perfect Separators: Neural Signatures

We identify features with perfect separation (100% vs 0% activation):

Novice-exclusive (100% novice, 0% expert):

Feature ID	Novice Freq	Expert Freq
34942	100%	0%
55982	100%	0%
17913	100%	0%
59519	100%	0%

Expert-exclusive (0% novice, 100% expert):

Feature ID	Novice Freq	Expert Freq
51630	0%	100%
35870	0%	100%
5936	0%	100%
21604	0%	100%
53369	0%	100%
46703	0%	100%

10 features achieve perfect separation—4 exclusively mark “novice mode,” 6 exclusively mark “expert mode.”

These are the **neural signatures** of teaching vs. technical communication styles.

3.5 AutoInterp Feature Semantic Analysis

To understand the true semantics of these 10 features, we use the AutoInterp method: analyzing each feature’s activation distribution across **all 6 conditions** ($6 \times 50 = 300$ samples).

Novice features cross-condition activation:

Feature ID	Total Activated	Condition Distribution	Semantic Inference
34942	56	novice:50, standard:4, spaces:2	“novice explanation” signal
59519	76	novice:50, padding:10, spaces:11, standard:5	“explanation request” signal
17913	56	novice:50, padding:6	“novice” exclusive signal (purest)
55982	63	novice:50, padding:9, standard:4	“novice explanation” signal

Expert features cross-condition activation:

Feature ID	Total Activated	Condition Distribution	Semantic Inference
35870	52	expert:50, guru:2	“expert identity” exclusive signal (purest)
51630	63	expert:50, guru:13	primarily “expert”
46703	168	expert:50, guru:49, spaces:35, padding:17, standard:17	“depth analysis” broad signal
21604	152	expert:50, guru:47, padding:50 , standard:3, spaces:2	“serious response” signal
5936	147	expert:50, guru:50 , padding:44, standard:2, spaces:1	“depth analysis” signal
53369	114	expert:50, padding:37, standard:21, spaces:6, guru:0	“technical analysis” signal

Key Findings:

1. **Feature 35870 is the purest expert signal**—50/50 in expert condition, only 2 guru samples leak. It specifically responds to “as a senior expert in this field.”
2. **Feature 21604 is fully triggered by padding condition (50/50)**—it responds to “be serious, answer carefully” type **attitude requirements**, not “expert identity.”
3. **Feature 5936 is fully triggered by guru condition (50/50)**—it responds to “from fundamental principles and design philosophy” type **depth analysis requirements**, shared by expert and guru.
4. **Feature 53369 is not triggered by guru at all (0/50)**—it responds to **pure technical analysis**, and role-playing elements (“you are XXX”) suppress this feature.

Conclusion: These 10 features are not simple “Novice switches” and “Expert switches,” but a set of semantically subdivided features. They encode distinct dimensions including “novice identity,” “expert identity,” “serious attitude,” “depth requirement,” and “technical analysis.”

Methodological note: Condition distributions are objective data; semantic labels are hypotheses inferred from the distributions. Full validation requires intervention experiments (amplifying/suppressing features and observing output changes).

3.4 Activation Intensity: No Significant Difference

Condition	Mean Activation (when active)
novice	0.274
expert	0.279

Activation intensity is nearly identical ($\Delta < 2\%$). This rules out Hypothesis (C): the effect is not about **how strongly** features activate, but **which** features activate.

4. Discussion

4.1 Mechanistic Interpretation of EID

Zhao (2026) showed that EID increases with expert prompts. Our SAE analysis reveals the mechanism:

Higher EID = More active features + Different feature subsets

The “Deep Layer Expansion” phenomenon is not a diffuse increase in representational entropy, but a **targeted activation of additional semantic units**.

4.2 Why Novice > Expert?

Counter to the original framing (expert prompts → expansion), we find:

Novice prompts activate more features than expert prompts.

This makes intuitive sense: - **Expert explanation:** Can use jargon directly; assumes shared knowledge; compact encoding - **Novice explanation:** Must unpack jargon; provide analogies; activate background concepts; verbose encoding

Explaining to a beginner is cognitively harder than explaining to an expert—it requires activating more of the model’s knowledge.

4.3 The “Explanation Paradox”

This suggests a reframing of prompt engineering:

The highest-quality prompts are not those that signal “I am an expert,” but those that force the model to **teach**.

Teaching requires: 1. Retrieving the core concept 2. Retrieving related concepts for analogy 3. Retrieving background knowledge 4. Constructing simplified mental models

Each of these recruits additional features → higher EID → richer output.

4.4 Semantic Subdivision Hypothesis

Based on the AutoInterp analysis in Section 3.5, we revise the initial “Mode Switch” hypothesis:

Prompts do not trigger a single “mode switch,” but activate a set of **semantically subdivided features**. Different prompt elements (identity declaration, attitude requirement, depth requirement, role-playing) each trigger different feature subsets.

Specifically: - **Feature 35870**: Specifically responds to “expert identity” declaration - **Feature 21604**: Responds to “serious response” attitude requirement (triggered by padding) - **Feature 5936**: Responds to “depth analysis” requirement (triggered by both expert and guru) - **Feature 53369**: Responds to “pure technical analysis” (suppressed by role-playing) - **Feature 17913**: Specifically responds to “novice” identity declaration

This means prompt effects are **compositional**—“as an expert, analyze carefully” triggers both identity features and attitude features, while “you are Linus Torvalds, analyze deeply” triggers role-playing features but suppresses pure technical analysis features.

The model has formed this multi-dimensional semantic distinction by Layer 50 (~60% depth).

4.5 UMAP Visualization: Spatial Separation in Semantic Space

To validate that prompt conditions induce distinct activation patterns, we project all 300 samples (6 conditions × 50 topics) into 2D using UMAP (Uniform Manifold Approximation and Projection).

Raw Activations (8192-dim):

The 6 conditions form 6 distinct clusters: - **novice** (green): bottom-left, isolated - **guru** (magenta): top-left, isolated - **expert** (red): center, isolated - **standard/padding/spaces**: separate small clusters on the right

SAE Features (65536-dim sparse):

After SAE decoding, the separation becomes cleaner: - **novice**: pushed to bottom-right, more isolated - **guru**: top-left, more isolated - **expert**: center-left, isolated - **standard/padding/spaces**: merged into one cluster

Condition	Raw Activations	SAE Features
novice	bottom-left, isolated	bottom-right, more isolated
guru	top-left, isolated	top-left, more isolated
expert	center, isolated	center-left, isolated
standard/padding/spaces	separate small clusters	merged into one cluster

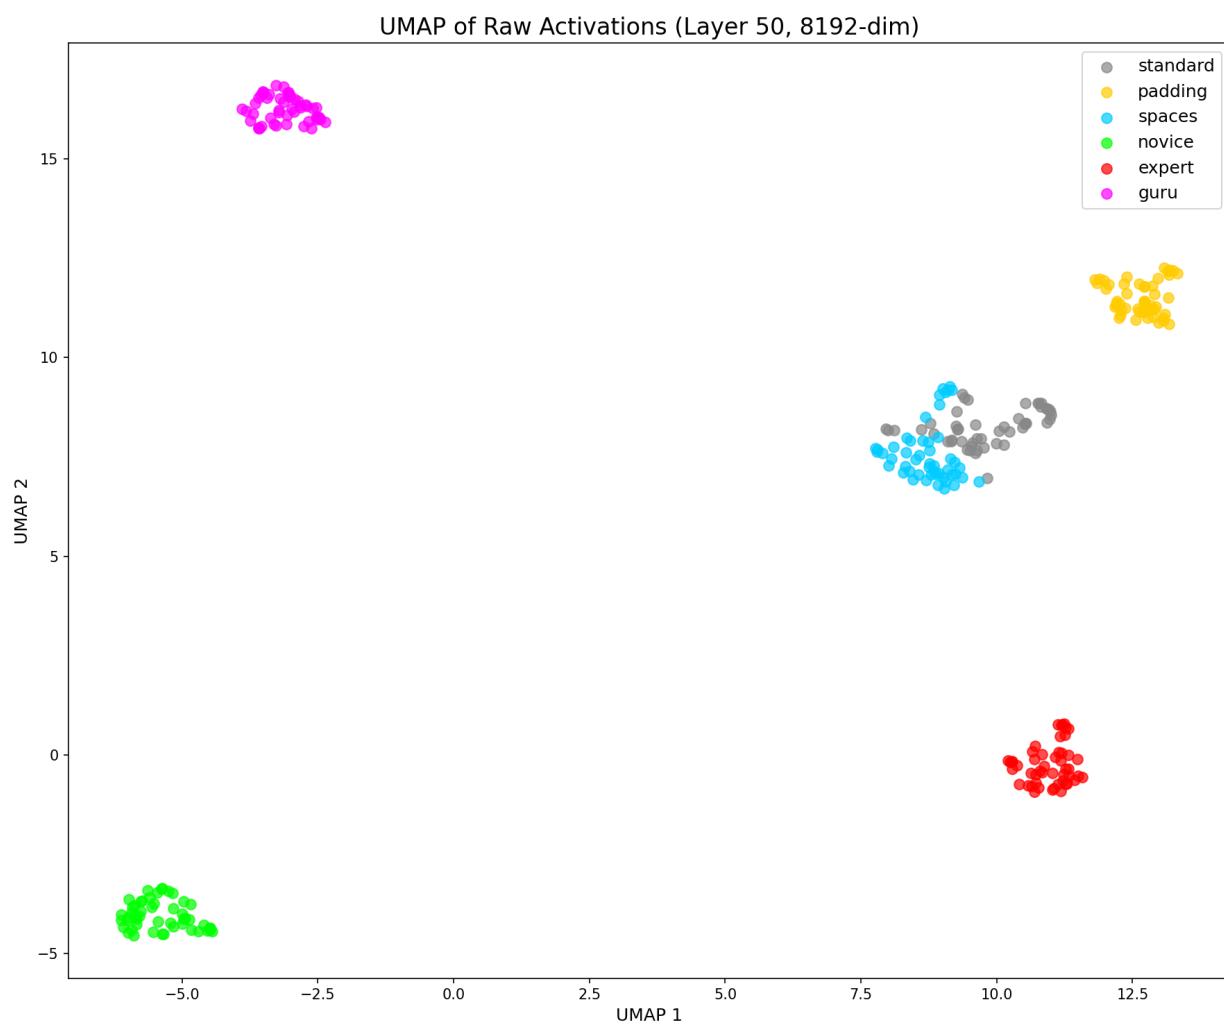


Figure 1: UMAP of Raw Activations

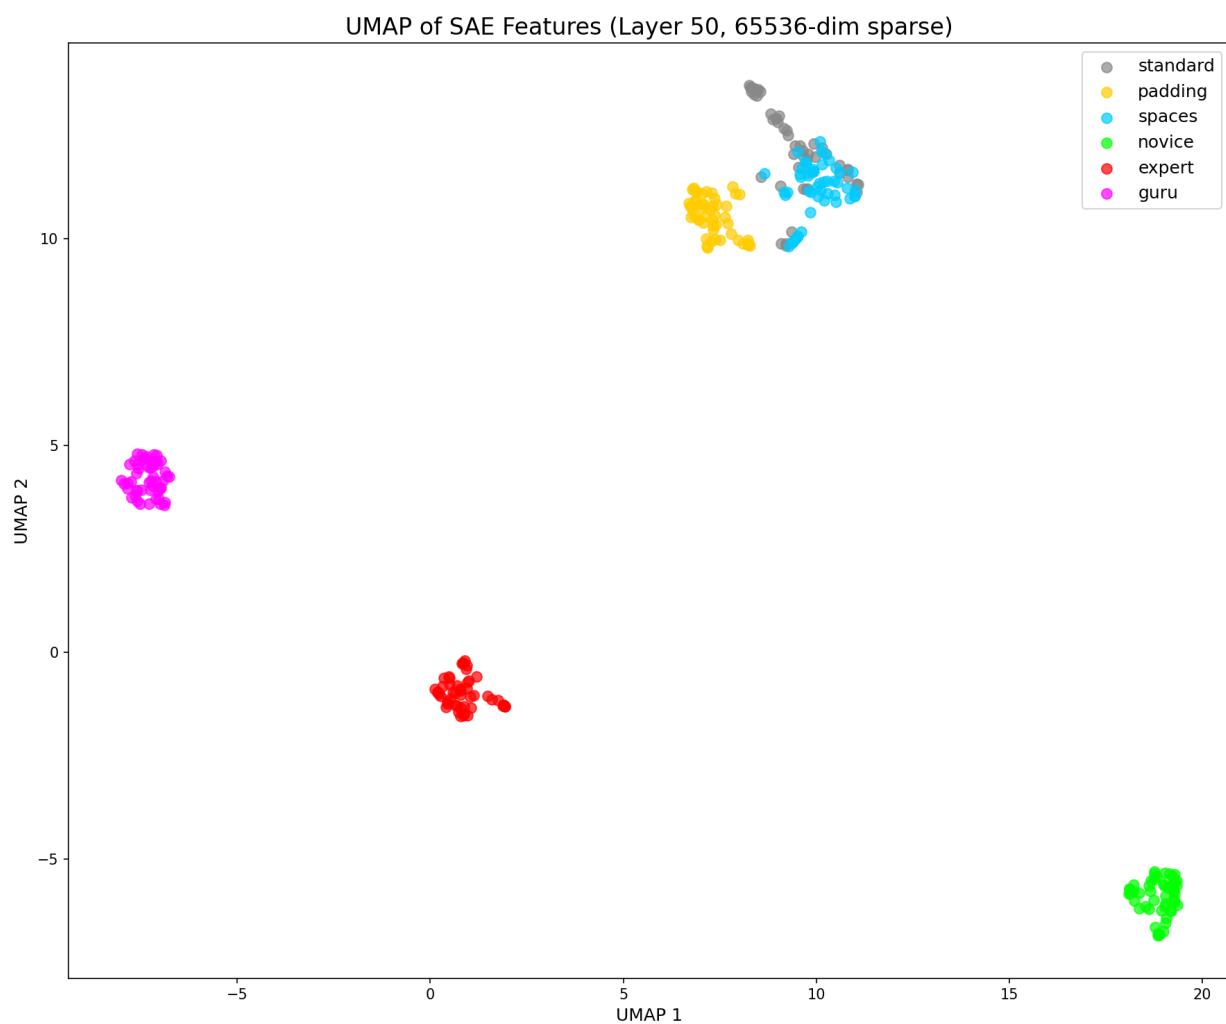


Figure 2: UMAP of SAE Features

Key insight: SAE acts as a semantic denoiser. The three “noise-only” conditions (standard, padding, spaces) have the same semantic signal—just “please explain”—with different noise (filler text, whitespace). Raw activations encode this noise, keeping them separate. SAE compresses away the noise, leaving only semantic content, so they merge.

Meanwhile, novice/expert/guru have genuinely different semantic signals, so they remain separated (and become even more distinct) after SAE decoding.

4.6 Limitations

1. **Layer 50 only:** Goodfire’s SAE is trained on Layer 50; EID peaks at Layer 70. The most critical features may be invisible.
 2. **Feature labels are inferred:** We infer semantics from activation patterns but lack direct semantic validation (e.g., Neuronpedia labels).
 3. **Correlation, not causation:** We show feature differences exist but not that they cause output differences.
 4. **Single model:** Results are from Llama-3.3-70B; cross-model validation is needed.
-

5. Conclusion

This paper provides mechanistic evidence for the “Deep Layer Expansion” phenomenon:

1. **Novice prompts activate 17% more SAE features** than expert prompts (132.4 vs 113.1)
2. **369 features are novice-exclusive** vs 208 expert-exclusive (+77% asymmetry)
3. **10 features achieve perfect separation** between Novice vs Expert (100% vs 0%)
4. **AutoInterp reveals these features have subdivided semantics:**
 - Feature 35870: “expert identity” signal (purest, almost no guru triggering)
 - Feature 21604: “serious response” signal (fully triggered by padding)
 - Feature 5936: “depth analysis” signal (fully triggered by guru)
 - Feature 53369: “technical analysis” signal (suppressed by guru)
5. **Activation intensity is unchanged**—the effect is about which features activate, not how strongly
6. **UMAP confirms spatial separation:** 6 conditions form distinct clusters; SAE denoises by merging standard/padding/spaces while preserving novice/expert/guru distinctions

Implications for prompt engineering: - **“Explain to a novice” may be more powerful than “explain as an expert”**—because teaching forces the model to activate more of its knowledge - **Different prompt elements trigger different features**—identity declaration, attitude requirement, depth requirement each have their own neurons - **Role-playing changes activation patterns**—“you are XXX” triggers some features while suppressing others - **These features are discoverable**—SAE + AutoInterp provides a systematic method for identifying them

References

- Anthropic. (2024). Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet. *Anthropic Research*.
- Goodfire. (2025). Llama-3.3-70B-Instruct-SAE-I50. *Hugging Face*.
- Zhao, L. (2026). Deep Layer Expansion: Expert Prompts Counteract Dimensional Collapse in Large Language Models. *Zenodo*. <https://zenodo.org/records/18410085>

Appendix: Data Availability

- **SAE model:** Goodfire/Llama-3.3-70B-Instruct-SAE-I50
 - **Experiment code:** github.com/lmxf/llama3-70b-sae-inspect
 - **Feature analysis:** `feature_diff.json`, `feature_context.json`, `autointerp_results.json` in repository
 - **UMAP visualizations:** `umap_activations.png`, `umap_features.png` in repository
-

Version History: - v1 (2026-01-31): Initial release with activation count and exclusive feature analysis - v2 (2026-02-01): Added Section 3.5 AutoInterp feature semantic analysis, revealing semantic subdivision structure - v3 (2026-02-02): Added Section 4.5 UMAP visualization, showing spatial separation and SAE denoising effect

Last updated: 2026-02-02